

Dismiss

Join GitHub today

GitHub is home to over 50 million developers working together to host and review code, manage projects, and build software together.


Sign up

Branch: master

Find file

Copy path

openshift-toolbox / ocp4_upi / docs / 0.prereqs.md

 **christianh814** Update 0.prereqs.md



123eb92 on Oct 17, 2019

1 contributor

Raw

Blame

History

552 lines (439 sloc) | 14.6 KB

Prereqs

This is probably the longest and most important part of this guide. The prereqs are "heavier" in this release. So I'll break it down into sections

- [Infra](#)
- [DHCP](#)
- [DNS](#)
- [Load Balancer](#)
- [WebServer](#)

Remember, a lot of this can be done for you if you use my [helper node](#) playbook.

Infra

This is a minimum setup

Create 6 "blank" VMs (or use BareMetal) for OpenShift itself (don't install an OS on this just yet)

- 3 Masters
 - 4 vCPUs
 - 16 GB RAM
 - 120 GB HD
- 2 Workers (I did 3; but 2 works)
 - 2 vCPUs
 - 8 GB RAM
 - 120 GB HD
- 1 Bootstrap (you remove this after the install is done)
 - 4 vCPUs
 - 16 GB RAM

- 120 GB HD

Create 1 VM for the LoadBalancer (I used CentOS 7 but any Linux will do)

- LoadBalancer VM
 - 2 vCPUs
 - 4 GB RAM
 - 30 GB HD

Create 1 VM for DNS (You can use Centos/Fedora/RHEL)

- DNS VM
 - 2 vCPUs (if using IDM use 4 vCPUs)
 - 4 GB RAM
 - 30 GB HD (if Using IDM set this to 50GB HD)
 - You can also use this for your [Webserver](#) too
 - You can install the [DHCP](#) components here too

If you don't have a lot of resources you can combine the LB and DNS server but I don't recommend it. If you don't have enough resources buy another server/box. See the [min requirements](#) for OpenShift 4...it's a BEAST!

DHCP

Right now, with 4.2, you can't set static ip addresses with the installer. So "static" DHCP (where you reserve the MAC address to a specific IP address in your network) is the way to go.

Here is a high level overview on how to do it...**I actually did mine on my router**; which is the way I recommend that you do it...but if you don't; here is how to install a simple DHCP server.

Install packages

```
yum -y install dhcp
```

Next set it up to where it'll start on boot

```
systemctl enable dhcpd
```

Make note of the MAC address of you LB and OpenShift infra. Use this to create

`/etc/dhcp/dhcpd.conf`. The below is an EXAMPLE. Use it to create your own...

```
authoritative;
ddns-update-style interim;
default-lease-time 14400;
max-lease-time 14400;

option routers                192.168.1.1;
option broadcast-address      192.168.1.255;
option subnet-mask            255.255.255.0;
option domain-name-servers   192.168.1.2, 192.168.1.3;
option domain-name            "example.com";

subnet 192.168.1.0 netmask 255.255.255.0 {
  pool {
    range 192.168.1.100 192.168.1.200;
    # Static entries
    host bootstrap { hardware ethernet 08:ed:b9:2d:82:4a; fixed-address
```

```

192.168.1.104; }
    host lb { hardware ethernet 08:ed:b9:2d:82:4b; fixed-address
192.168.1.105; }
    host master0 { hardware ethernet 01:1d:b9:2d:82:4c; fixed-address
192.168.1.100; }
    host master1 { hardware ethernet 02:ad:b9:2d:82:4d; fixed-address
192.168.1.101; }
    host master2 { hardware ethernet 03:4d:b9:2d:82:4e; fixed-address
192.168.1.102; }
    host worker0 { hardware ethernet 04:7d:b9:2d:82:4f; fixed-address
192.168.1.110; }
    host worker1 { hardware ethernet 05:bd:b9:2d:82:5e; fixed-address
192.168.1.111; }
    host worker2 { hardware ethernet 08:8d:b9:2d:82:1b; fixed-address
192.168.1.112; }
    # this will not give out addresses to hosts not listed above
    deny unknown-clients;
}
}

```

Configure firewall stuff

```

firewall-cmd --add-service=dhcp --permanent
firewall-cmd --reload

```

Now start the service

```
systemctl start dhcpd
```

^ Again....**I DON'T RECOMMEND THIS** ...if you can do this in your router...do it there (that's what I did).

DNS

OCP4 uses SRV records for autodiscovery. More information can be found [here](#). I used idm (a.k.a FreeIPA) for DNS. But that's a little heavy (it also installs Kerberos, LDAP, Dogtag, etc).

So if you JUST need DNS; here's a quick overview on how to set it up. (the below assumes Fedora but can work on RHEL/CentOS)

Install it

```
dnf -y install bind bind-utils
```

Edit the `/etc/named.conf` file to look similar to what's below. I made only 3 changes

- `listen-on` was changed to read `any`
- `allow-query` was changed to read `any`
- `forward only`; I added this for DNS to forward queries
- `forwarders { 8.8.8.8; 8.8.4.4; };` forwarding those queries to Google's DNS servers
- The entries for my zonefiles (noted by comments)

THE DEFAULT FILE `/etc/named.conf` **is fine** ...just modify/add the things I mentioned above

```

//
// named.conf
//
// Provided by Red Hat bind package to configure the ISC BIND named(8) DNS
// server as a caching only nameserver (as a localhost DNS resolver only).

```

```
//
// See /usr/share/doc/bind*/sample/ for example named configuration files.
//

options {
    listen-on port 53 { any; };
    listen-on-v6 port 53 { ::1; };
    directory      "/var/named";
    dump-file       "/var/named/data/cache_dump.db";
    statistics-file "/var/named/data/named_stats.txt";
    memstatistics-file "/var/named/data/named_mem_stats.txt";
    allow-query     { any; };

    /*
     * - If you are building an AUTHORITATIVE DNS server, do NOT enable recursion.
     * - If you are building a RECURSIVE (caching) DNS server, you need to enable
     *   recursion.
     * - If your recursive DNS server has a public IP address, you MUST enable access
     *   control to limit queries to your legitimate users. Failing to do so will
     *   cause your server to become part of large scale DNS amplification
     *   attacks. Implementing BCP38 within your network would greatly
     *   reduce such attack surface
     */
    recursion yes;

    /* Forwarders */
    forward only;
    forwarders { 8.8.8.8; 8.8.4.4; };

    dnssec-enable yes;
    dnssec-validation yes;

    managed-keys-directory "/var/named/dynamic";

    pid-file "/run/named/named.pid";
    session-keyfile "/run/named/session.key";

    /* https://fedoraproject.org/wiki/Changes/CryptoPolicy */
    include "/etc/crypto-policies/back-ends/bind.config";
};

logging {
    channel default_debug {
        file "data/named.run";
        severity dynamic;
    };
};

zone "." IN {
    type hint;
    file "named.ca";
};

##### Add what's between these comments #####
zone "example.com" IN {
    type master;
    file "example.com";
};

zone "1.168.192.in-addr.arpa" IN {
    type master;
    file "192.168.1.db";
};
#####

include "/etc/named.rfc1912.zones";
include "/etc/named.root.key";
```

Create `/var/named/example.com` zone file with the following entries. Making sure you've set your SRV records. NOTE, whatever your subdomain is (in this case it's `ocp4`) **WILL BE** the name of your cluster when you install.

```
$TTL 1W
@      IN      SOA      ns1.example.com.      root (
                                2019052300      ; serial
                                3H              ; refresh (3 hours)
                                30M             ; retry (30 minutes)
                                2W              ; expiry (2 weeks)
                                1W )           ; minimum (1 week)
                                IN      NS      ns1.example.com.
                                IN      MX 10    smtp.example.com.
;
;
ns1     IN      A        192.168.1.2
smtp    IN      A        192.168.1.2
;
; The api points to the IP of your load balancer
api.ocp4      IN      A        192.168.1.105
api-int.ocp4  IN      A        192.168.1.105
;
; The wildcard also points to the load balancer
*.apps.ocp4   IN      A        192.168.1.105
;
; Create entry for the bootstrap host
bootstrap.ocp4 IN      A        192.168.1.104
;
; Create entries for the master hosts
master0.ocp4  IN      A        192.168.1.100
master1.ocp4  IN      A        192.168.1.101
master2.ocp4  IN      A        192.168.1.102
;
; Create entries for the worker hosts
worker0.ocp4  IN      A        192.168.1.110
worker1.ocp4  IN      A        192.168.1.111
worker2.ocp4  IN      A        192.168.1.112
;
; The ETCD cluster lives on the masters...so point these to the IP of the masters
etcd-0.ocp4   IN      A        192.168.1.100
etcd-1.ocp4   IN      A        192.168.1.101
etcd-2.ocp4   IN      A        192.168.1.102
;
; The SRV records are IMPORTANT...make sure you get these right...note the trailing dot
at the end...
_etcd-server-ssl._tcp.ocp4      IN      SRV      0 10 2380 etcd-0.ocp4.example.com.
_etcd-server-ssl._tcp.ocp4      IN      SRV      0 10 2380 etcd-1.ocp4.example.com.
_etcd-server-ssl._tcp.ocp4      IN      SRV      0 10 2380 etcd-2.ocp4.example.com.
;
;EOF
```

Alright, now create the `/var/named/192.168.1.db` file for reverse DNS. This is important because this is how the host sets it's hostname on boot.

```
$TTL 1W
@      IN      SOA      ns1.example.com.      root (
                                2019052300      ; serial
                                3H              ; refresh (3 hours)
                                30M             ; retry (30 minutes)
                                2W              ; expiry (2 weeks)
                                1W )           ; minimum (1 week)
                                IN      NS      ns1.example.com.
;
; syntax is "last octet" and the host must have fqdn with trailing dot
100     IN      PTR      master0.ocp4.example.com.
```

```
101      IN      PTR      master1.ocp4.example.com.
102      IN      PTR      master2.ocp4.example.com.
;
104      IN      PTR      bootstrap.ocp4.example.com.
;
105      IN      PTR      api.ocp4.example.com.
105      IN      PTR      api-int.ocp4.example.com.
;
110      IN      PTR      worker0.ocp4.example.com.
111      IN      PTR      worker1.ocp4.example.com.
112      IN      PTR      worker2.ocp4.example.com.
;
;
;EOF
```

Make sure SELinux is happy

```
restorecon -vR /var/named
```

If you're using a firewall; open up those ports

```
firewall-cmd --permanent --add-port=53/tcp --add-port=53/udp
firewall-cmd --add-port=53/tcp --add-port=53/udp
```

Enable/start it

```
systemctl enable --now named
```

Make sure it's up

```
systemctl status named
```

LB

I used HAProxy on CentOS 7 ...it's pretty straight forward.

Install it on your lb (192.168.1.105) server

```
yum -y install haproxy
```

Edit the `/etc/haproxy/haproxy.cfg` file to look like this...

```
#-----
# Example configuration for a possible web application.  See the
# full configuration options online.
#
#   http://haproxy.1wt.eu/download/1.4/doc/configuration.txt
#
#-----

#-----
# Global settings
#-----
global
    # to have these messages end up in /var/log/haproxy.log you will
    # need to:
    #
```

```

# 1) configure syslog to accept network log events. This is done
#   by adding the '-r' option to the SYSLOGD_OPTIONS in
#   /etc/sysconfig/syslog
#
# 2) configure local2 events to go to the /var/log/haproxy.log
#   file. A line like the following can be added to
#   /etc/sysconfig/syslog
#
#   local2.*                                /var/log/haproxy.log
#
log                127.0.0.1 local2

chroot             /var/lib/haproxy
pidfile            /var/run/haproxy.pid
maxconn            4000
user               haproxy
group              haproxy
daemon

# turn on stats unix socket
stats socket /var/lib/haproxy/stats

#-----
# common defaults that all the 'listen' and 'backend' sections will
# use if not designated in their block
#-----
defaults
    mode                http
    log                  global
    option               httplog
    option               dontlognull
    option http-server-close
    option forwardfor    except 127.0.0.0/8
    option               redispatch
    retries              3
    timeout http-request 10s
    timeout queue        1m
    timeout connect      10s
    timeout client       1m
    timeout server       1m
    timeout http-keep-alive 10s
    timeout check        10s
    maxconn              3000

#-----

listen stats
    bind :9000
    mode http
    stats enable
    stats uri /
    monitor-uri /healthz

frontend openshift-api-server
    bind *:6443
    default_backend openshift-api-server
    mode tcp
    option tcplog

backend openshift-api-server
    balance source
    mode tcp
    server boot 192.168.1.104:6443 check
    server master-0 192.168.1.100:6443 check
    server master-1 192.168.1.101:6443 check
    server master-2 192.168.1.102:6443 check

```

```
frontend machine-config-server
    bind *:22623
    default_backend machine-config-server
    mode tcp
    option tcplog

backend machine-config-server
    balance source
    mode tcp
    server boot 192.168.1.104:22623 check
    server master-0 192.168.1.100:22623 check
    server master-1 192.168.1.101:22623 check
    server master-2 192.168.1.102:22623 check

frontend ingress-http
    bind *:80
    default_backend ingress-http
    mode tcp
    option tcplog

backend ingress-http
    balance source
    mode tcp
    server worker-0 192.168.1.110:80 check
    server worker-1 192.168.1.111:80 check
    server worker-2 192.168.1.112:80 check

frontend ingress-https
    bind *:443
    default_backend ingress-https
    mode tcp
    option tcplog

backend ingress-https
    balance source
    mode tcp
    server worker-0 192.168.1.110:443 check
    server worker-1 192.168.1.111:443 check
    server worker-2 192.168.1.112:443 check

#-----
```

Set up the firewall

```
firewall-cmd --permanent --add-port=9000/tcp \
--add-port=443/tcp \
--add-port=80/tcp \
--add-port=6443/tcp \
--add-port=6443/udp \
--add-port=22623/tcp \
--add-port=22623/udp
firewall-cmd --reload
```

Make SELinux Happy

```
setsebool -P haproxy_connect_any on
```

Start/enable HAProxy

```
systemctl enable haproxy --now
```

Verify


```
systemctl status haproxy
```

Visit the status page

```
firefox http://192.168.1.105:9000
```

^ This page is important to keep up when installing

Webserver

You need a place to store the bios file and ignition files. You can create a webserver for this (make sure you change `/etc/httpd/conf/httpd.conf` file to "Listen" on port `8080` since the HAProxy LB is listening on `80`)

```
yum -y install httpd
systemctl enable --now httpd
firewall-cmd --add-service=8080/tcp --permanent
firewall-cmd --reload
```

^ I didn't do this...since I'm on a linux laptop I just ran the following from my laptop

```
python3 -m http.server -d /path/to/install/dir 8080
```

But the apache webserver is fine too

Conclusion

Do the following commands to check everything

DNS

etcd SRV records

```
$ dig _etcd-server-ssl._tcp.ocp4.example.com SRV +short
0 10 2380 etcd-0.ocp4.example.com.
0 10 2380 etcd-2.ocp4.example.com.
0 10 2380 etcd-1.ocp4.example.com.
```

Also check your LB

```
$ dig api.ocp4.example.com +short
192.168.1.105
$ dig api-int.ocp4.example.com +short
192.168.1.105
```

You should check the hosts names too of your nodes

LB

Simple curl will return the health of the LB itself

```
$ curl api.ocp4.example.com:9000/healthz
<html><body><h1>200 OK</h1>
```

```
Service ready.  
</body></html>
```

[return to the index page](#)